

Artificial Intelligence¹

Roman Barták

Department of Theoretical Computer Science and Mathematical Logic

Problem Solving: Uninformed Search

Example

- Intelligent agents are supposed to maximize their performance measure. This can be simplified if the agent can adopt a **goal** and aim at satisfying it.

A model problem and a possible way to solve it

- an agent is in the city of Arad (Romania)
- the performance measure contains many factors (improve its suntan, take in the sights, enjoy the nightlife, ...), which make decisions hard
- now, suppose the agent has a non-refundable ticket to fly out of Bucharest the following day
- goal formulation** – getting to Bucharest – greatly simplifies the agent's decision problem
- a goal is a set of world states, where the goal is satisfied, but we also need to assume other world states and actions for moving between – **problem formulation**
 - states could correspond to being in major towns (the exact location is too fine)
 - actions could correspond to driving from one major town to another (acceleration and breaking are too fine)
- How to find a path to Bucharest, if three roads lead out from Arad, one toward Sibiu, one to Timisoara, and one to Zerind?
 - Suppose the agent has a map of Romania and can explore possible journeys (**search**), select the best one, and then finally **execute** the actions.




- Simple reflex agent** „only” transfers the current percept to one action.
- Goal-based agent** can plan action sequences that achieve agent's goals.

- We will now explore one kind of goal-based agent – a **problem-solving agent**.
- What is a **problem** and what is its **solution**?
- How to **find a solution**?
- Search** algorithms
BFS, DFS, DLS, IDS, BiS

Note:

First, we will explore “uninformed” search algorithms, that is, algorithms that do not exploit any problem specific information (and use an atomic representation).



Problem solving

Problem solving consists of four steps:

- goal formulation**
 - What are the desired states of the world?
- problem formulation**
 - What are the states and actions assumed to reach the goal?
- problem solving**
 - How to find the best sequence of states reaching the goal?
- solution execution**
 - If we know the actions, how to execute them?




```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
          state, some description of the current world state
          goal, a goal, initially null
          problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```

Problem formulation

Well-defined problems consist of:

- **the initial state**
 - $in(Arad)$
- a description of possible **actions** can be done via a **transition model** describing for each state possible actions and states to which the action leads
 - $SUCCESSOR-FN(in(Arad)) = \{go(Sibiu), in(Sibiu)\}$
 - implicitly defines the **state space** (the set of all states reachable from the initial state by any sequence of actions)
 - a **path** is a sequence of states connected by actions.
- **the goal test**
 - a function determining whether a given state is a goal state or not
 - $\{in(Bucharest)\}$
- **a path cost**
 - a function that assigns a numeric cost to each path (reflects the performance measure)
 - We assume that the cost of a path can be described as a sum of the costs of individual actions along the path.



A **solution to a problem** is an action sequence that leads from the initial state to a goal state.

An **optimal solution** is a solution that has the lowest path cost among all solutions.

Abstraction

- In the problem formulation, we used
 - **abstraction of world states**
 - we ignored weather, traffic conditions, ...
 - **abstraction of actions**
 - we ignored turning on the radio, looking out of the window, ...
- **Abstraction** is the process of removing detail from representation.
- What is the **appropriate level of abstraction**?
 - the abstraction is **valid**
 - we can expand any abstract solution into a solution in the more detailed world
 - we can find a way from any place in Arad to any place in Sibiu
 - the abstraction is **useful**
 - carrying out each of the actions in the solution is easier than the original problem
 - the path from Arad to Sibiu can be executed by average driving agent


Intelligent agents would be completely swamped by the real world without using abstractions.

Toy problems

- **Toy problems** are used to illustrate and comparison of solving techniques.


The vacuum world

states (robot \times dirt)
start, goal
actions (L, R, S)




The 8-puzzle

states (block position)
start, goal
actions (L, R, U, D)



The 8-queens

- place queens on chessboard without attacks
- incremental formulation (adding queens)
 - complete-state formulation (moving queens)




Real problems

- That is what matters!

Route-finding problems

states (places)
start (here and now), goal (there, on-time)
successor
cost function (time, cost,...)




Touring problems

goal (to visit each place at least once)
states (current and visited places)
TSP (Travelling Salesman Problem), NP-hard

Product assembly problems


state (hand location \times components)
goal (assembled product)
successor (movement of „hinges“)
cost (total assembly time)
protein design from a sequence of amino acids



- Having formulated the problem, how do we solve it?


State space search

- start in the **initial state** (root node)
- check** whether the **initial state** is the **goal state**
- if the state is not a goal state then **expand the state**, it generates a set of new states
- select a next state** using a **search strategy**




```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

- State space is different from search tree – world state is different from search node.**



- Node** of the search tree consists of:

- a **current state**
- a link to its **parent**
- action** leading from parent to the current state
- cost** of path from the root $g(n)$
- depth** (the number of steps from the root)



- Fringe (frontier)** is a set of nodes not yet expanded.



- Nodes from the **fringe** are called **leaves**.
- The algorithms can be described using the following operations over the fringe represented as a **queue**:
 - MAKE-QUEUE(element,...)
 - EMPTY?(queue)
 - FIRST(queue)
 - REMOVE-FIRST(queue)
 - INSERT(element, queue)
 - INSERT-ALL(elements, queue)



```
function TREE-SEARCH(problem, fringe) returns a solution, or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
    fringe ← INSERT-ALL(EXPAND(node, problem), fringe)
```

```
function EXPAND(node, problem) returns a set of nodes
  successors ← the empty set
  for each action, result in SUCCESSOR-FN[problem](STATE[node]) do
    s ← a new NODE
    PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
    PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s] ← DEPTH[node] + 1
    add s to successors
  return successors
```

We can evaluate an **algorithm's performance** in four ways:

- **completeness**
 - Is the algorithm guaranteed to find a solution when there is one?
- **optimality**
 - Does the strategy find the optimal solution?
- **time complexity**
 - How long does it take to find a solution?
- **space complexity**
 - How much memory is needed to perform search?
- Time and space complexity are always considered with respect to some measure of problem difficulty.
 - **branching factor b** (maximum number of successors of any node) – useful for implicit representation of world states (the initial state and successors)
 - **depth d** (the path length from the root to a shallowest goal node)
 - **path length m** (the maximum length of any path in the state space)
- **search cost** (how much time and space we need to find a solution)
- **total cost** (combines the search cost and the path cost of the solution found)




- Today we will cover **uninformed (blind) search**

- **no additional information about states** beyond that provided in the problem definition
- it can generate successors and distinguish a goal state from a non-goal state




- Next lesson will cover **informed (heuristic) search**

- information to distinguish “**more promising**” non-goal states from other states
- for example by estimating distance to some goal state



Breadth-first search

- **All nodes are expanded at a given depth** in the search tree before any nodes at the next level are expanded.



Breadth-first search (analysis)

- It is a **complete method** (provided the branching factor is finite).
- The shallowest goal state is not necessarily the optimal one!
 - **BFS is optimal if the path cost is a non-decreasing function of the depth of the node**
- **Time complexity** (the number of visited nodes for the goal at depth d)
 - $1 + b + b^2 + b^3 + \dots + b^d + b(b^d - 1) = O(b^{d+1})$
- **Space complexity**
 - store every expanded node in the explored set
 - **$O(b^{d+1})$**

The memory requirements are a bigger problem than is the execution time.

$b = 10$
10,000 nodes/sec.
1000 bytes/node

depth	nodes	time	memory
2	1100	0.11 sec.	1 megabyte
4	111100	11 seconds	106 megabytes
6	10^7	19 minutes	10 gigabytes
8	10^9	31 hours	1 terabyte
10	10^{11}	129 days	101 terabytes
12	10^{13}	35 years	10 petabytes
14	10^{15}	3523 years	1 exabyte

Uniform-cost search

- Modifying BFS for **finding an optimal solution**.
- Expand the node n with the **lowest path cost** $g(n)$.
- Beware of zero-cost steps – they can cause cycling!
- **Completeness** (and optimality) can be guaranteed if the cost of each step is lower-bounded by some ϵ .
- **Time (and space) complexity**
 - depends on the path cost rather than on the depth
 - $O(b^{1+\lceil C^*/\epsilon \rceil})$, where C^* is the cost of the optimal solution
 - **Could be much worse than for BFS**, because the algorithm prefers long paths with “cheap” steps over shorter paths with more expensive steps.



Depth-first search (analysis)

- If the algorithm selects a wrong path, it may not find a goal during tree search (it is **not complete**) and of course it may **not find optimum**.
- **Time complexity**
 - $O(b^m)$, where m is the maximum depth
 - it may happen that $d \ll m$, where d is the depth of the optimum
- **Space complexity**
 - needs to store only a single path from the root to a leaf node, along with the remaining unexpanded sibling nodes
 - expanded node can be removed as soon as all its descendants have been fully explored
 - $O(bm)$, where m is the maximum visited depth
 - Space complexity can be decreased via **backtracking!**
 - generate one successor (instead of all) $\rightarrow O(m)$ states
 - modify the current state rather than copying it (we must be able to undo each modification when going back to generate the next successor $\rightarrow O(1)$ states and $O(m)$ actions

Depth-first search

- Always **expand the deepest node**
- **until reaching the level, where the nodes have no successors and then the search “backs up” to the node at the shallower level**
- achieved by using a **LIFO** (last-in-first-out) queue for the frontier
- Frequently implemented in a recursive way.



Depth-limited search

- The failure of DFS in infinite state spaces can be alleviated by supplying DFS with a **predetermined depth limit l** .
 - nodes at depth l are treated as if they have no successors
 - terminates with a solution, with a failure, or with cut-off
 - **time complexity $O(b^l)$, space complexity $O(bl)$**
 - if $l < d$, then the algorithm is not complete (d is a depth of solution)
 - if $d \ll l$, then the algorithm explores many useless nodes

How to decide the depth limit?

```

function DEPTH-LIMITED-SEARCH( problem, limit) returns soln/fail/cutoff
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)
function RECURSIVE-DLS( node, problem, limit) returns soln/fail/cutoff
  cutoff-occurred?  $\leftarrow$  false
  if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
  else if DEPTH[node] = limit then return cutoff
  else for each successor in EXPAND(node, problem) do
    result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
    if result = cutoff then cutoff-occurred?  $\leftarrow$  true
    else if result  $\neq$  failure then return result
  if cutoff-occurred? then return cutoff else return failure
    
```

- using knowledge of the problem
- for example, for path finding for 20 cities, we can use the depth limit 19
- if we study the map carefully, we can discover that any city can be reached from any other city in at most 9 steps (the **diameter** of the state space)

Iterative deepening

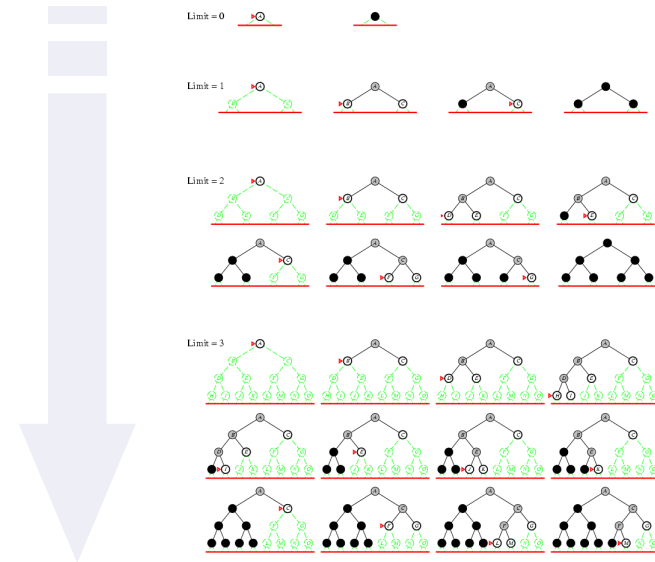
How to make the depth limited search complete?

- by **gradually increasing the limit**
- combines the benefits of BFS and DFS
 - **completeness** (when the branching factor is finite)
 - **optimal** (when the path cost is a non-decreasing function of the depth of the node)
 - low **memory** consumption $O(bd)$
 - What about **time complexity**?
 - $d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d = O(b^d)$
 - in fact, this is even better than BFS, which explores one more level
 - Ex. ($b=10, d=5$): BFS = 1.111.100, DLS = 111.110, IDS = 123.450

Iterative deepening is the preferred uninformed search method when the search space is large and the depth of the solution is not known.


```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  inputs: problem, a problem
  for depth ← 0 to ∞ do
    result ← DEPTH-LIMITED-SEARCH(problem, depth)
    if result ≠ cutoff then return result
```

Iterative deepening (example)



Bidirectional search

We can run two simultaneous searches – **one forward from the initial state and the other backward from the goal** (hoping that the two searches meet in the middle).



Why?

- $b^{d/2} + b^{d/2} \ll b^d$
- Ex. ($b=10, d=6$): BFS = 11.111.100, BiS = 22.200

How?

- Before expanding a node, verify that it is not present in the frontier of the other search tree.
- One of the frontiers must be kept in memory $O(b^{d/2})$ so the intersection check can be done. With a hash table it will take constant time.
- If using breadth-first search, we obtain a **complete algorithm** (still the first solution found may not be optimal and some additional search is necessary to make sure there isn't any shortcut).

Backward search?

- Forward search explores the set of states, but **what is explored during backward search and how?**
 - If the goal is a single state, we can go through the states (path Arad → Bucharest).
 - If the goal is a set of goal states given explicitly (vacuum world), we can use a **dummy state** such that all the goal states are its predecessors or we can use a **meta-state** to describe a set of states.
 - The worst case is when there is an **abstract description of the goal state** (8-queens).
- Another problem is **definition of predecessors** of a state.
 - The predecessors of n are all those states that have n as a successor.
 - Again, the most complex case is when the state transition is defined using a general function.




Criterion	Breadth-First	Uniform-cost	Depth-First	Depth-limited	Iterative deepening	Bidirectional search
Complete?	YES*	YES*	NO	YES, if $l \geq d$	YES*	YES*
Time	b^{d+1}	$b^{C^*/\epsilon}$	b^m	b^l	b^d	$b^{d/2}$
Space	b^{d+1}	$b^{C^*/\epsilon}$	bm	bl	bd	$b^{d/2}$
Optimal?	YES*	YES*	NO	NO	YES*	YES*

- b – the branching factor
- d – the depth of the shallowest solution
- C^* – the cost of optimal solution
- ϵ – the minimal step cost (the minimal action cost)
- m – the maximum depth of the search tree
- l – the depth limit
- * complete if b is finite (BFS)
- optimal if step costs are all identical

- So far we ignored one of typical problems of tree search – **expansion of already visited states**.

- sometimes, it is not a problem (a good problem formulation, ex. 8-queens)
- sometimes, the search tree is much bigger than the search space



• How to resolve it?

```

function GRAPH-SEARCH( problem, fringe ) returns a solution, or failure
  closed ← an empty set
  fringe ← INSERT( MAKE-NODE( INITIAL-STATE[problem] ), fringe )
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT( fringe )
    if GOAL-TEST[problem]( STATE[node] ) then return SOLUTION( node )
    if STATE[node] is not in closed then
      add STATE[node] to closed
      fringe ← INSERTALL( EXPAND( node, problem ), fringe )

```

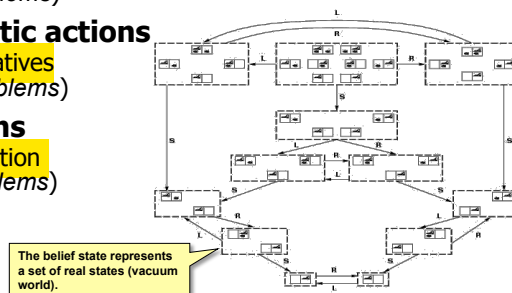
- remember already expanded states – **closed states**

- if the closed state is visited again, it is treated as it has no successors

- So far we assumed static, fully observable, discrete and deterministic environment.

- What if some **information is missing**?

- **no sensors** (the agent does not know the current state)
 - work with **belief states** – sets of real states – looking for a belief state containing only goal states (*conformant problems*)
- **non-deterministic actions**
 - plans with alternatives (*contingency problems*)
- **unknown actions**
 - solved by exploration (*exploration problems*)



© 2013 Roman Barták

Department of Theoretical Computer Science and Mathematical Logic
bartak@ktiml.mff.cuni.cz