

# Artificial Intelligence<sup>1</sup>

Roman Barták

Department of Theoretical Computer Science and Mathematical Logic

Problem Solving: Local and On-line Search Techniques

- We know how to use **heuristics in search**
  - BFS, A\*, IDA\*, RBFS, SMA\*
- **Today:**
  - **What if the path is not important?**
    - Local search: HC, SA, BS, GA
  - **What if the world changes?**
    - on-line search, LRTA\*




## Local search

- So far we systematically explored all paths possibly going to the goal and the path itself was a part of the solution.
- For some problems (e.g. 8-queens) the **path is not relevant to the problem, only the goal** is important.
- For such problems we can try **local search** techniques.
  - we also keep a single state only (constant memory)
  - in each step we slightly modify the state
  - usually, the path is not stored
  - the method can also look for an **optimal state**, where the optimality is defined by an objective function (defined for states).
    - For example, for the 8-queens problem the objective function can be defined by the number of conflicting queens– this is extra information about the quality of states.



## Local search - terminology

- Local search can be seen as a move in the state-space **landscape**, where coordinates define the state and **elevation corresponds to the objective function.**




## Hill climbing

- From the neighborhood the algorithm **selects the state with the best value of the objective function** and moves there (**hill climbing**).
  - knows only the neighborhood
  - the current state is forgotten

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
inputs: problem, a problem
local variables: current, a node
                 neighbor, a node
current ← MAKE-NODE(INITIAL-STATE[problem])
loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```


• # conflicts = 17  
• state change = change row of a queen  
• random selection among more best neighbours



## Hill climbing - problems

HC is a **greedy algorithm** – goes to the best neighbour with looking ahead

- local optimum** (the state such that each neighbour is not better)
  - HC cannot escape local optimum
- ridges** (a sequence of local optima)
  - difficult for greedy algorithms to navigate
- plateaux** (a flat area of the state-space landscape)
  - shoulder – progress is still possible
  - HC may not find a solution (cycling)




## Hill climbing - versions

- stochastic HC**
  - chooses at random from among the uphill moves; the probability of selection can vary with the steepness of the uphill move
  - usually converges more slowly than steepest ascent
  - in some landscapes, it finds better solutions
- first-choice HC**
  - implements stochastic HC until a successor better than the current state is generated
  - a good strategy when a state has many (thousands) of successors
- random-restart HC**
  - conducts a series of HC searches from randomly generated initial states (restart)
  - can escape from a local optimum
  - if HC has a probability  $p$  of success then the expected number of restarts required is  $1/p$
  - a very efficient method for the N-queens problem ( $p \approx 0.14$  i.e. 7 iterations to find a goal)

## Simulated annealing

- HC never makes the "downhill" moves towards states with lower value so the algorithm is not complete (can get stuck on a local optimum).
- Random walk** – that is moving to a successor chosen randomly – is complete but extremely inefficient.
- Simulated annealing** combines hill climbing with random walk
  - motivation in metallurgy – process to harden metals by heating them to a high temperature and then gradually cooling them (allowing the material to reach a low-energy crystalline state)
  - the algorithm picks a random move and accepts it if:
    - it improves the situation
    - it worsens the situation but this is allowed with a probability given by some temperature value and how much the state worsens; the temperature is decreasing according to a cooling scheme

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
inputs: problem, a problem
       schedule, a mapping from time to "temperature"
local variables: current, a node
                 next, a node
                 T, a "temperature" controlling prob. of downward steps
current ← MAKE-NODE(INITIAL-STATE[problem])
for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\text{next}] - \text{VALUE}[\text{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E / T}$ 
```



Keeping just one node in memory might seem to be an extreme reaction to the problem of memory limitations.

Can we exploit available memory better?

### Local beam search algorithm

- keeps track of  $k$  states rather than just one
- at each step, all the successors of all  $k$  states are generated
  - if any one is a goal, the algorithm halts
- otherwise, it selects the  $k$  best successors and repeats



### This is not running $k$ restarts of HC in parallel!

- useful information is passed among the parallel search threads
- The algorithm quickly abandons unfruitful searches and moves its resources to where the most is being made
- can suffer from a lack of diversity
  - stochastic beam search helps alleviate this problem ( $k$  successors are chosen at random with the probability being an increasing function of state value)
  - resemble the process of natural selection



- A variant of stochastic beam search in which successors are generated by combining two parent states (sexual reproduction)

- Begin with a set of  $k$  randomly generated states – **population**
  - Each state is represented as a string over a finite alphabet (DNA)
  - **fitness** function evaluates the states (objective function)
- Select a pair of states for reproduction (probability of selection is given by the fitness function)
- For each pair choose a **crossover** point from the positions in the string
- Combine **offsprings** to a new state
- Each location is subject to random **mutation** with a small probability



**function** GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual  
**inputs:** *population*, a set of individuals  
 FITNESS-FN, a function that measures the fitness of an individual

**repeat**

*new\_population*  $\leftarrow$  empty set

**loop for**  $i$  **from** 1 **to** SIZE(*population*) **do**

$x \leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

$y \leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

*child*  $\leftarrow$  REPRODUCE( $x, y$ )

**if** (small random probability) **then** *child*  $\leftarrow$  MUTATE(*child*)

**add** *child* **to** *new\_population*

*population*  $\leftarrow$  *new\_population*

**until** some individual is fit enough, or enough time has elapsed

**return** the best individual in *population*, according to FITNESS-FN

**function** REPRODUCE( $x, y$ ) **returns** an individual

**inputs:**  $x, y$ , parent individuals

$n \leftarrow$  LENGTH( $x$ )


$c \leftarrow$  random number from 1 to  $n$

**return** APPEND(SUBSTRING( $x, 1, c$ ), SUBSTRING( $y, c + 1, n$ ))

- So far we have concentrated on **offline search**
  - compute a complete solution
  - then execute the solution without assuming percepts
- **Online search** is different in
  - **interleave computation and action**
    - select an action
    - execute an action
    - observe the environment
    - compute the next action
  - a good idea **in dynamic and semidynamic environments**
  - helpful in nondeterministic domains (unknown actions and unknown results of actions)



- Online search is useful **for agents executing actions** (useless for pure computation).
- The agent knows only the following **information**:
  - Actions(s)** – a list of actions allowed in state  $s$
  - $c(s,a,s')$**  – the step cost function (cannot be used until the agent knows state  $s'$ )
  - Goal-Test(s)** – identifying the goal state
- We assume the following:
  - agent **can recognize a visited state**
  - (agent can build a world map)
  - actions are deterministic
  - agent has an **admissible heuristic**  $h(s)$



Quality of online algorithms can be measured by **comparing with the offline solution** (knowing the best solution in advance).

- Competitive ratio**


= quality of online solution / quality of the best solution

- can be  $\infty$ , for example for a **dead-end state** (if some actions are **irreversible**).

**Claim:** No algorithm can avoid dead ends in all state spaces.

Proof (adversary argument)

Agent has visited states  $S$  and  $A$  must make the same decision in both, but in one situation, the agent reaches dead-end.



Assume that the **state space is safely explorable** (some goal state is reachable from every reachable state).

- No bounded competitive ratio can be guaranteed if there are paths of unbounded cost.
- Adversary argument can be used to arbitrarily extend any path.



Hence it is common to describe the performance of online search algorithms in **terms of the size of the entire state space** rather than just the depth of the shallowest goal.

- Opposite to offline algorithms such as A\* online algorithms **can discover successors only** for a node that the agent physically occupies.
- It seems better to **expand nodes in a local order** as done for example by DFS.

```

function ONLINE-DFS-AGENT( $s'$ ) returns an action
inputs:  $s'$ , a percept that identifies the current state
static: result, a table, indexed by action and state, initially empty
unexplored, a table that lists, for each visited state, the actions not yet tried
unbacktracked, a table that lists, for each visited state, the backtracks not yet tried
 $s$ ,  $a$ , the previous state and action, initially null


if GOAL-TEST( $s'$ ) then return stop
if  $s'$  is a new state then  $unexplored[s'] \leftarrow ACTIONS(s')$ 
if  $s$  is not null then do
   $result[a, s] \leftarrow s'$ 
  add  $s$  to the front of  $unbacktracked[s']$ 
if  $unexplored[s']$  is empty then
  if  $unbacktracked[s']$  is empty then return stop
  else  $a \leftarrow$  an action  $b$  such that  $result[b, s'] = POP(unbacktracked[s'])$ 
else  $a \leftarrow POP(unexplored[s'])$ 
 $s \leftarrow s'$ 
return  $a$ 

```

learns outcome of actions

The state can be visited more times in a single journey (it is leaved to different states) so we need to remember where to go back – **Ariadna thread**

to go back we need a reverse action to the actions used to reach the state




stav	unEX	unBT	rUP	rDN	rLF	rRG
(1,a)	{}	(2,a)	(2,a)	-	-	(1,b)
(2,a)	{}	(1,a)	-	(1,a)	-	-
(1,b)	LF, RG	(1,a)	(2,b)	-		
(2,b)	DW	(1,b)	(3,b)		-	-
(3,b)	DW	(3,a), (2,b)	-		(3,a)	-
(3,a)		(3,b)	-	-	-	(3,b)

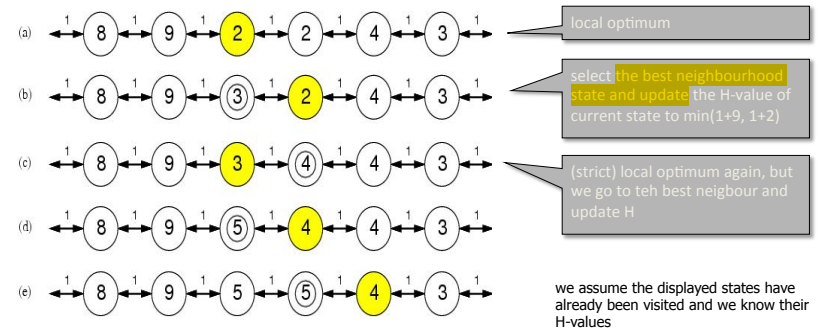
- In the **worst case**, every link is traversed exactly twice (forward and backward).
- This is optimal for exploration, but for finding a goal, the agent's competitive ratio could be arbitrarily bad.
  - an online variant of iterative deepening solves this problem
- On-line DFS works only in state spaces where actions are reversible.**

- **Hill climbing is an on-line algorithm.**

- keeps a **single state** (the current physical state)
- does **local steps** to the neighbouring states
- in its simplest form **cannot escape local optima**
  - Beware! **Restarts cannot be used in online search!**
  - We can still use **random work**.
    - A random walk will **eventually find a goal** or complete its exploration, provided that the space is finite.
    - The process **can be very slow**.
    - In the following example, a **random walk will take** exponentially many steps to find the goal (at each step, backward progress is twice as likely as forward progress).



- We can exploit available memory to remember visited states and hence leave local optima.
  - **H(s)** – the current best estimate of path length from **s** to the goal (equals  $h(s)$  at the beginning)



- Algorithm **LRTA\*** makes local steps and learns the result of each action as well as a better estimate of distance to the goal ( $H$ ).

