

Artificial Intelligence¹

Roman Barták

Department of Theoretical Computer Science and Mathematical Logic

Adversarial Search: Games


Games

- **Mathematical game theory** (a branch of economics) views any multi-agent environment as a game, provided that the impact of each agent on others is significant.
 - environments with many agents are called economies (rather than games)
- AI deals mainly with **turn-taking, two-player zero-sum** games (one player wins, the other one loses).
 - **deterministic** games vs. stochastic games
 - **perfect** information vs. imperfect information
 - Why games in AI? Because games are:
 - hard to play
 - easy to model (not that many actions)
 - funny



	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon, monopoly
imperfect information		bridge, poker, scrabble, nuclear war


- So far we assumed a single-agent environment, but what if there are **more agents and some of them „playing“ against us?**
- Today we will discuss **adversarial search** a.k.a. **game playing**, as an example of a **competitive multi-agent environment**.
 - **deterministic, turn-taking, two-player zero-sum games of perfect information** (tic-tac-toe, chess)
 - optimal (perfect) decisions (minimax, alpha-beta)
 - imperfect decisions (cutting off search)
 - stochastic games (backgammon)



Problem setting

- We consider two players **MAX** and **MIN**
 - **MAX moves first**, and then the players take turns moving until the game is over
 - we are looking for the strategy of MAX
- Again, we shall see game playing as a search problem:
 - **initial state**: specifies **how the game is set up at the start**
 - **successor function**: results of the moves (move, state)
 - the **initial state** and the **successor function** define the **game tree**
 - **terminal test**: true, when the game is over (a goal state)
 - **utility function**: final numeric value for a game that ends in terminal state (win, loss, draw with values +1, 0, -1)
 - **higher values are better for MAX, while lower values are better for MIN**

- Two players place X and O in an empty square until a line of three identical symbols is reached or all squares are full.



- Classical search is looking a **(shortest) path to a goal state**.
- Search for games is looking for a **path to the terminal state with the highest utility**, but MIN has something to say about it.
- MAX is looking for a contingent strategy, which specifies
 - MAX's move in the initial state
 - MAX's moves in the states resulting from every possible response by MIN
 - an **optimal strategy** leads to outcomes at least as good as any other strategy when one is playing an infallible opponent

- The **optimal strategy** can be determined from the **minimax value** of each node computed as follows:

MINIMAX-VALUE(n) =

UTILITY(n)

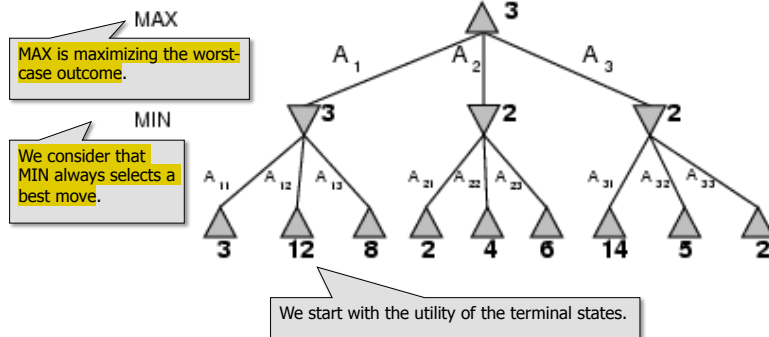
if n is a terminal state

$\max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$

if MAX plays in n

$\min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$

if MIN plays in n



```
function MINIMAX-DECISION(state) returns an action
     $v \leftarrow \text{MAX-VALUE}(\textit{state})$ 
    return the action in  $\text{SUCCESSORS}(\textit{state})$  with value  $v$ 
```

```
function MAX-VALUE(state) returns a utility value
    if  $\text{TERMINAL-TEST}(\textit{state})$  then return  $\text{UTILITY}(\textit{state})$ 
     $v \leftarrow -\infty$ 
    for  $a, s$  in  $\text{SUCCESSORS}(\textit{state})$  do
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$ 
    return  $v$ 
```


The algorithm assumes that the player plays optimally. Otherwise, the utility is even higher

```
function MIN-VALUE(state) returns a utility value
    if  $\text{TERMINAL-TEST}(\textit{state})$  then return  $\text{UTILITY}(\textit{state})$ 
     $v \leftarrow \infty$ 
    for  $a, s$  in  $\text{SUCCESSORS}(\textit{state})$  do
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$ 
    return  $v$ 
```

• Time complexity $O(b^m)$
 • Space complexity $O(bm)$
 (b - #actions in states, m - #moves)

Minimax for more players


- For multiplayer games we can use a **vector of utility values** – this vector gives the utility of the state from each player's viewpoint.



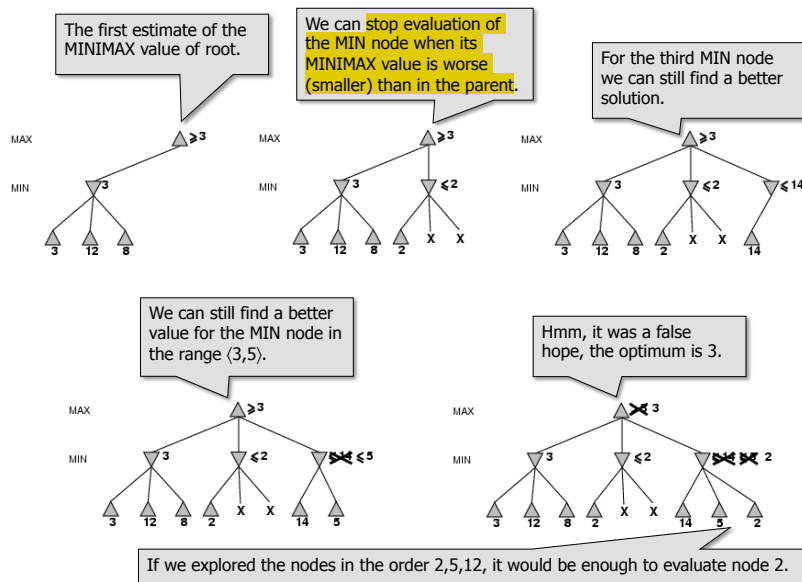
- Multiplayer games usually involve **alliances**, whether formal or informal, among the players.
 - Alliances seems to be a natural consequence of optimal strategies for each player.
 - For example, suppose A and B are in weak positions and C is in a stronger position. Then it is often optimal for both A and B to attack C rather than each other. Of course, as soon as C weakens under the joint onslaught, the alliance loses its value.

Improving minimax

- The minimax algorithm always finds an optimal strategy, but it has to explore a complete game tree.
- Can we speed-up the algorithm?
 - YES! We do not need to explore all states, if they are "very bad".
 - α - β pruning** eliminates branches that cannot possibly influence the final decisions.



α - β pruning - example

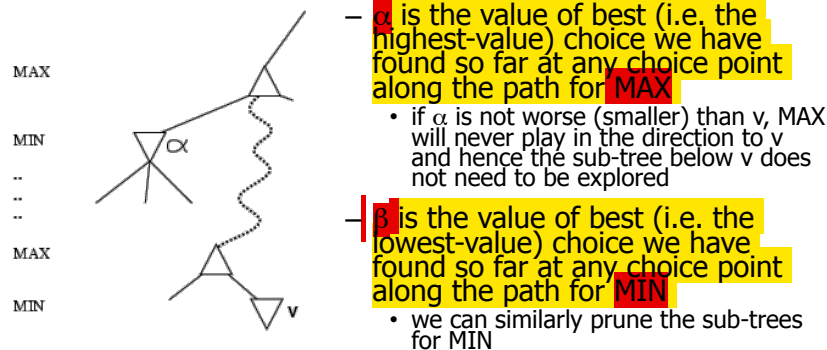


Algorithm α - β

```
function ALPHA-BETA-SEARCH(state) returns an action
  inputs: state, current state in game
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in  $\text{SUCCESSORS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
   $\alpha$ , the value of the best alternative for MAX along the path to state
   $\beta$ , the value of the best alternative for MIN along the path to state
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow -\infty$ 
  for  $a, s$  in  $\text{SUCCESSORS}(\text{state})$  do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  inputs: state, current state in game
   $\alpha$ , the value of the best alternative for MAX along the path to state
   $\beta$ , the value of the best alternative for MIN along the path to state
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow +\infty$ 
  for  $a, s$  in  $\text{SUCCESSORS}(\text{state})$  do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```



Properties:

- By cutting off the sub-trees we do not miss optimum.
- By "perfect ordering" we can decrease time complexity to $O(b^{m/2})$, which gives a branching factor \sqrt{b} (b for minimax), so we can solve a tree roughly twice as deep as minimax in the same amount of time.

Evaluation function

- Returns an estimate of the expected utility of the game from a given position (similar to the heuristic function h).
- Obviously, quality of the algorithm depends on the quality of evaluation function.

Properties:

- terminal states must be ordered in the same way as if ordered by the true utility function
- the computation must not take too long
- for nonterminal states, the evaluation function should be strongly correlated with the actual chances of winning
 - given the limited amount of computation, the best the algorithm can do is make a guess about the final outcome
- How to construct such a function?



- Both minimax and α - β have to **search all the way to terminal states**.
 - This is not practical for bigger depths (depth = #moves to reach a terminal state).
- We can **cut off search** earlier and apply a heuristic evaluation function to states in the search.
 - does not guarantee finding an optimal solution, but
 - can finish search in a given time
- **Realisation:**
 - terminal test \rightarrow cutoff test
 - utility function \rightarrow heuristic evaluation function EVAL

Evaluation function - examples

Expected value

- based on selected features of states, we can define various *categories* (equivalence classes) of states
- each category is evaluated based on the proportion of winning and losing states
 - $EVAL = (0.72 \times +1) + (0.20 \times -1) + (0.08 \times 0) = 0.52$

„Material“ value

- estimate the numerical contribution of each feature
 - chess: pawn = 1, knight = bishop = 3, rook = 5, queen = 9
- combine the contributions (e.g. weighted sum)
 - $EVAL(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$
 - The sum assumes independence of features!
 - It is possible to use non-linear combination.



White moves first and Black wins

- The situation may change dramatically by assuming one more move after the cut-off limit.

Identical material value (better for Black) for both states, but **White wins the right position** by capturing the queen.




- quiescent**: if the opponent can capture a chess-man then the estimate is not stable and it is better to explore a few more moves (for example only selected moves)

horizon effect


- the unavoidable bad situation can be delayed after the cut-off limit (horizon) and hence it is not recognized as a bad state

Black has a better material value, but if White changes a pawn to a queen, then **White wins**. Black may consider checking the white king so the situation does not look so bad.



- In real life, many unpredictable external events can put us into unforeseen situations.
- Games mirror **unpredictability** by including a random element, such as throwing of dice.

Backgammon



- the goal is to move all one's pieces off the board (clockwise)
- who finishes first, wins
- dice are rolled to determine the legal moves
 - the total travelled distance

There are four legal moves for White: (5-10,5-11), (5-11,19-24), (5-10,10-16), (5-11,11-16)

Singular extension

- explore the sequence of moves that are "clearly better" than all other moves
- a fast way to explore the area after the depth limit (quiescent is a special case)

Forward pruning


- some moves at a given state are not assumed at all (a human approach)
- dangerous as it can miss the optimal strategy
- safe, if symmetric moves are pruned

Transposition tables

- similarly to classical search, we can remember already evaluated states for the case when they are reached again by a different sequence of moves

- Game tree is extended with **chance nodes** (in addition to MAX and MIN nodes) describing all rolls of dice.
 - 36 results for two dice, 21 without symmetries (5-6 and 6-5)
 - chance for double is 1/36, other results 1/18

Chance nodes are added to each layer, where the move is influenced by randomness. MAX rolls the dice here.



- Instead of the MINIMAX value, we use **expected MINIMAX value** (based on probability of chance actions):

$$\text{EXPECTIMINIMAX-VALUE}(n) =$$

$\text{UTILITY}(n)$

$\max_{s \in \text{successors}(n)} \text{EXPECTIMINIMAX-VALUE}(s)$

$\min_{s \in \text{successors}(n)} \text{EXPECTIMINIMAX-VALUE}(s)$

$\sum_{s \in \text{successors}(n)} P(s) \cdot \text{EXPECTIMINIMAX}(s)$

if n is a terminal node

if MAX plays in n

if MIN plays in n


if n is a chance node



Beware of the evaluation function (for cut-off)

- the absolute value of nodes may play a role
- the values should be a linear transformation of expected utility in the node

The left tree is better for A_1 while the right tree is better for A_2 , though the order of nodes is identical.




Time complexity $O(b^{mn^m})$, where n is the number of random moves

- it is not realistic to reach a bigger depth especially for larger random branching

Using cut-off à la α - β

- we can cut-off the chance nodes if the evaluation function is bounded
- the expected value can be bounded when the value is not yet computed



Example: how to become rich (a different view of cards)

- **Situation 1:** Trail A leads to a gold pile while trail B leads to a road-fork. Go left and there is a mound of diamonds, but go right and a bus will kill you (diamonds are more valuable than gold). Where to go?
 - the best choice is B and left
- **Situation 2:** Trail A leads to a gold pile while trail B leads to a road-fork. Go right and there is a mound of diamonds, but go left and a bus will kill you. Where to go?
 - B a right
- **Situation 3:** Trail A leads to a gold pile while trail B leads to a road-fork. Select the correct side and you will reach a mound of diamonds, but select a wrong side and a bus will kill you. Where to go?
 - a reasonable agent (not risking the death;-) goes A
- This is the same case as in the previous slide. We do not know what happens at the road-fork B. In the card game, we do not know which card ($\heartsuit 4$ or $\diamondsuit 4$) the opponent has, 50% chance of failure.
- **Lesson learnt:** We need to assume information that we will have at a given state (the problem of using $\clubsuit 9$ is that MAX plays differently when all cards are visible).



- Card games may look like the stochastic games, but the dice are rolled just once at the beginning!
- Card games are an example of games with **partial observability** (we do not see opponent's cards).

Example: card game "higher takes" with open cards

Situation 1: MAX: $\heartsuit 6 \diamondsuit 6 \clubsuit 9 \spadesuit 8$ MIN: $\heartsuit 4 \spadesuit 2 \clubsuit 10 \spadesuit 5$

1. MAX gives $\clubsuit 9$, MIN confirms colour $\clubsuit 10$ MIN wins
 2. MIN gives $\spadesuit 2$, MAX gives $\diamondsuit 6$ MIN wins
 3. MAX gives $\heartsuit 6$, MIN confirms colour $\heartsuit 4$ MAX wins
 4. MIN gives $\clubsuit 5$, MAX confirms colour $\clubsuit 8$ MAX wins
- $\clubsuit 9$ is the optimal first move for MAX

Situation 2: MAX: $\heartsuit 6 \diamondsuit 6 \clubsuit 9 \spadesuit 8$ MIN: $\heartsuit 4 \spadesuit 2 \clubsuit 10 \spadesuit 5$

- a symmetric case, $\clubsuit 9$ is again the optimal first move for MAX

Situation 3: MIN hides the first card ($\heartsuit 4$ or $\diamondsuit 4$), what is the optimal first move for MAX now?

- Independently of $\heartsuit 4$ and $\diamondsuit 4$ the optimal first move was $\clubsuit 9$, so it is the first optimal move now too.
- **Really?**



Chees

- 1997 Deep Blue wins over Kasparov 3.5 – 2.5
- 2006 „regular“ PC (DEEP FRITZ) beats Kramnik 4 – 2

Checkers

- 1994 Chinook became the official world champion
- 29. 4. 2007 solved – optimal policy leads to draw

Go

- branching factor 361 makes it challenging
- today computers play at a master level (using Monte Carlo methods based on the UCT scheme)

Bridge

- 2000 GIB was twelve at world championship
- Jack and Wbridge5 play at the level of best players

