

Introduction to Machine Learning

NPFL 054

<http://ufal.mff.cuni.cz/course/npfl054>

Barbora Hladká
hladka@ufal.mff.cuni.cz

Martin Holub
holub@ufal.mff.cuni.cz

Charles University in Prague,
Faculty of Mathematics and Physics,
Institute of Formal and Applied Linguistics

Lecture 11 – Feature selection

- **Curse of dimensionality**
- **Feature selection heuristics**
- **Bayes error**

Curse of dimensionality

According to the Wikipedia:

The **curse of dimensionality** refers to various phenomena that arise when analyzing and organizing data **in high-dimensional spaces** (often with hundreds or thousands of dimensions) **that do not occur in low-dimensional settings**.

The common theme of these problems is that **when the dimensionality increases, the volume of the space increases so fast that the available data become sparse**. This sparsity is problematic for any method that requires statistical significance. In order to obtain a statistically sound and reliable result, the **amount of data needed** to support the result often **grows exponentially with the dimensionality**.

Also organizing and searching data often relies on detecting areas where objects form groups with similar properties; **in high dimensional data however all objects appear to be sparse and dissimilar** in many ways which prevents common data organization strategies from being efficient.

Curse of dimensionality

High dimensional data is difficult to work because there are not enough observations to get good/reliable statistical estimates

Consider a simple example. Random vector of binary variables with the same binomial distributions. (X_1, X_2, \dots, X_n) .

- Observe the frequency of different vector values if e.g.
 $\Pr(X_i = 1) = 1/2$ or $\Pr(X_i = 1) = 1/10$.

- If $\Pr(X_i = 1) = 1/10$, then $\Pr(1, 1, \dots, 1) = 1/10^n$ (!)

Thus, the need for data grows exponentially with the number of features!

Curse of dimensionality

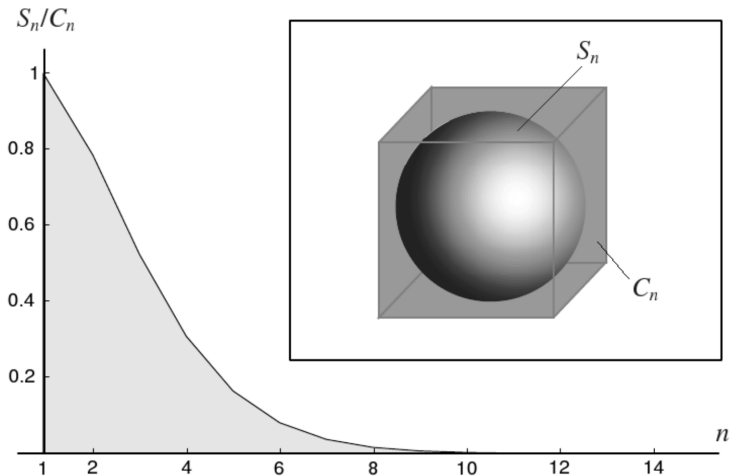
High-dimensional data is difficult to work not only because there are not enough observations to get good estimates... but also because **data distributed in a high dimensional space necessarily tends to be very sparse!**

This fact implies long distances between randomly distributed points

Consider a simple example. Uniformly distributed random points in an n -dimensional hypercube. – What will be their average/expected distance from the origin?

Curse of dimensionality – a geometrical illustration

Source: “The curse of dimensionality” by Mario Köppen



Ratio of the volumes of unit hypersphere and embedding hypercube

Curse of dimensionality

... also, in high-dimensional spaces there are long distances between randomly distributed points ...

Another example with uniformly distributed random points in an n -dimensional hypercube:

- What will be the mutual distance between two randomly selected points?

Curse of dimensionality

... also, in high-dimensional spaces there are long distances between randomly distributed points ...

Another example with uniformly distributed random points in an n -dimensional hypercube:

- What will be the mutual distance between two randomly selected points?

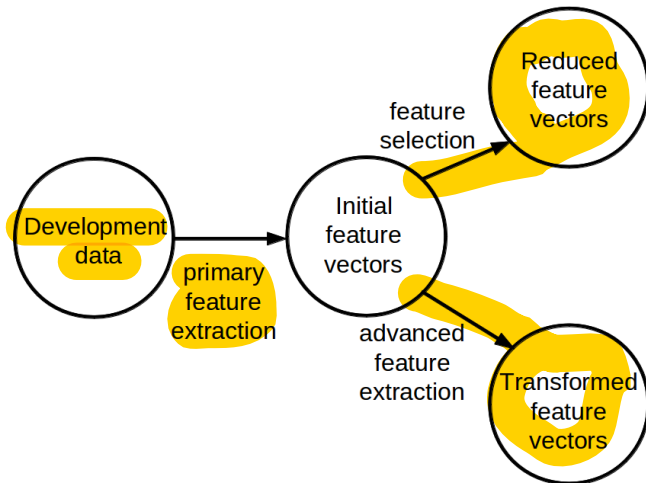
“Near neighbours” often do not exist!

– Instead, typically you have only many “far neighbours” ...

... and you cannot recognize the “similar ones”

Feature extraction and feature selection

Processes and terminology related to feature extraction/selection



Introduction to practical feature selection

Goal of the feature selection process = find a **minimum set of variables that contain all the substantial information** about predicting the target value

Introduction to practical feature selection

Goal of the feature selection process = find a minimum set of variables that contain all the substantial information about predicting the target value

- reduced feature space dimension in the dataset
- enhanced generalization and improved prediction performance by reducing overfitting (irrelevant input features may lead to overfitting)
- better chance to analyse the impact/importance of the features
- removing highly dependent features (some learning methods do not work well with them)
- lower model complexity and improved model interpretability
- feasible/shorter training times

Feature selection methods

Practical feature selection methods are heuristic

Practical feature selection methods are heuristic

Feature selection methods can be basically divided into

- **filters** – select feature subsets as a pre-processing step, independently of the learning method
- **wrappers** – use a machine learning algorithm in conjunction with internal cross validation procedure to score feature subsets by measuring their predictive power
- **embedded methods** – perform feature selection during the process of training

Simple methods in R: the FSelector package

```
> packageDescription('FSelector')
```

Description

This package provides functions for selecting attributes from a given dataset. Attribute subset selection is the process of identifying and removing as much of the irrelevant and redundant information as possible.

Feature ranking

~ aka variable importance metrics/measures

- We need a (real) function to evaluate how useful a feature is
- Frequently/mostly used:
Information Gain, Gini Index, Chi-square, correlation coefficient, etc.
 - see Wikipedia: “Feature Selection”
 - see the FSelector package in R
- Disadvantages: such methods consider only one variable's contribution without other variables' influences
- However, using them you can easily recognize
 - really useful ones
 - completely unuseful ones
 - highly dependent/correlated ones

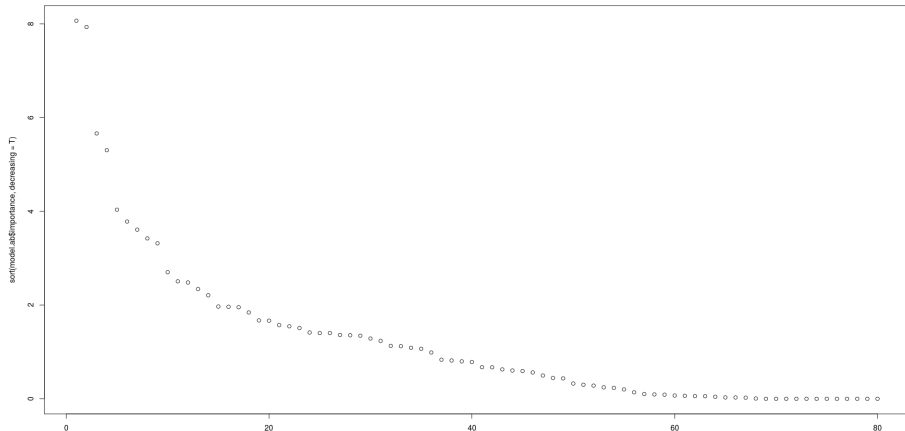
Practical methods for feature selection

Selected examples

- **Filters and wrappers**
 - greedy forward selection
 - greedy backward elimination
- **Variable importance produced by ensembles**
- **Feature selection by Lasso**
- **SVM-RFE – Recursive Feature Elimination**

Variable importance (AdaBoost) – *cry*

Example of the variable importance distribution



SVM-RFE feature selection algorithm

Example of succesfully combined heuristics

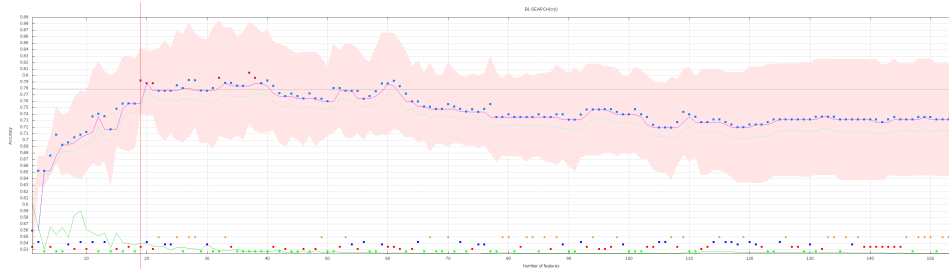
Algorithm 2 Recursive feature elimination using the SVM learner with cross-validated optimization of the SVM parameter *cost* in each iteration step.

Input: Training data set and the initial feature set

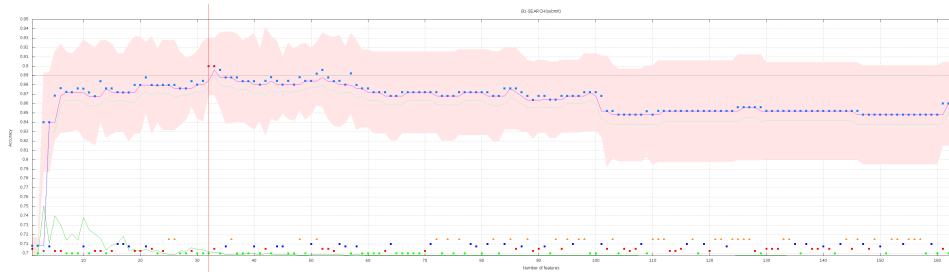
Output: The best SVM classifier M_{\max} and the corresponding feature subset S_{\max}

- 1: $K \leftarrow$ the initial feature set size
 - 2: $S_K \leftarrow$ the initial feature set
 - 3: **for** $k \leftarrow K$ **downto** 1 **do**
 - 4: learn a linear SVM model using the feature set S_k and tune its parameter cost
 - 5: $M_k \leftarrow$ the best tuned linear SVM model using the feature set S_k
 - 6: $f_{\text{worst}} \leftarrow$ the least useful feature in the model M_k
 - 7: $S_{k-1} \leftarrow S_k \setminus \{f_{\text{worst}}\}$
 - 8: **end for**
 - 9: $M_{\max} \leftarrow$ choose the best model from $\{M_i\}_{i=1}^K$
 - 10: $S_{\max} \leftarrow$ the best feature subset corresponding to the best model M_{\max}
-

SVM-RFE – *cry*

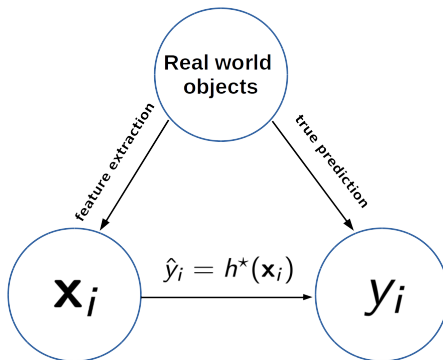


SVM-RFE – *submit*



Bayes classifier and Bayes error

Imagine that you are able to develop a really optimal classifier.
Is the zero test error always feasible?



Bayes classifier and Bayes error

Imagine that you are able to develop a really optimal classifier.
Is the zero test error always feasible?

The **Bayes classifier** minimises the probability of misclassification

Thus, by definition, error produced by the Bayes classifier is irreducible and is called *Bayes error*.

What is the lowest possible error rate

Bayes classifier assigns each example to the most likely class, given its feature values

$$\hat{y} = \max_y \Pr(y | \mathbf{x})$$

The Bayes classifier produces the lowest possible test error rate, so called **Bayes error rate**

$$1 - E (\max_y \Pr(y | \mathbf{x}))$$

What is the lowest possible error rate

Practical view on your development data

Are there identical feature vectors in your data set?

- Get the same feature vectors
- How many of them have the same target value?